

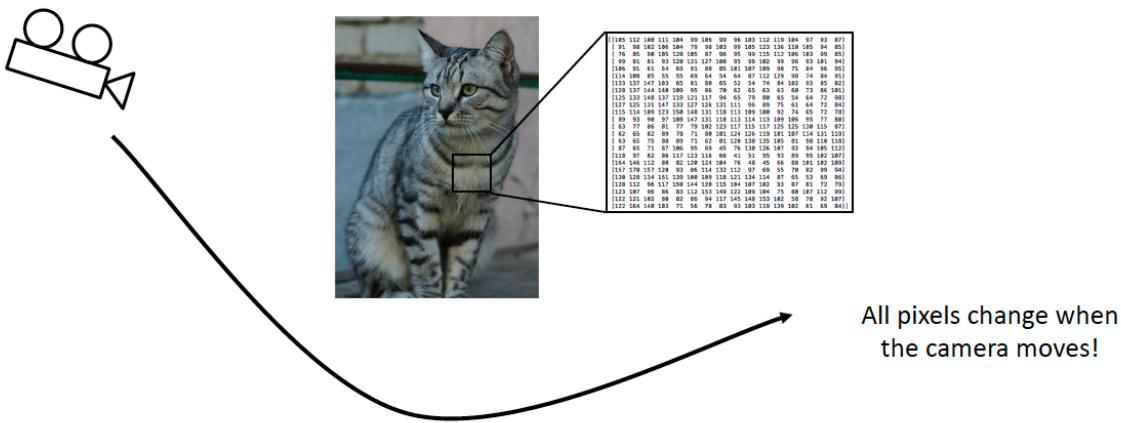
Lecture 2 - Image Classification

--> A core computer vision task

Input: image --> **Output:** Assign image to one of a fixed set of categories

1. Challenges (due to semantic gap)

- Viewpoint Variation



- Intraclass Variation



- Fine-Grained Categories



- **Background Clutter**
- **Illumination Changes**
- **Deformation**
- **Occusion 遮挡**

2. Image Classifier

- **Data-Driven Approach**
 - S1. Collect a dataset of images and labels
 - S2. Use Machine Learning to train a classifier
 - S3. Evaluate the classifier on new images
- **Classifier**

```
def train(images, labels):
    # Machine learning!
    return model
```

————→ Memorize all data and labels

```
def predict(model, test_images):
    # Use model to predict labels
    return test_labels
```

————→ Predict the label of the most similar training image

 - **Nearest Neighbor**
 - Distance metric to compare images (p: each pixel)

$$\text{L1 distance: } d_1(I_1, I_2) = \sum_p |I_1^p - I_2^p|$$

test image				training image				pixel-wise absolute value differences			
56	32	10	18	10	20	24	17	46	12	14	1
90	23	128	133	8	10	89	100	82	13	39	33
24	26	178	200	12	16	178	170	12	10	0	30
2	0	255	220	4	32	233	112	2	32	22	108

add → 456

```

import numpy as np

class NearestNeighbor:
    def __init__(self):
        pass

    def train(self, X, y):
        """ X is N x D where each row is an example. Y is 1-dimension of size N """
        # the nearest neighbor classifier simply remembers all the training data
        self.Xtr = X
        self.ytr = y

    def predict(self, X):
        """ X is N x D where each row is an example we wish to predict label for """
        num_test = X.shape[0]
        # lets make sure that the output type matches the input type
        Ypred = np.zeros(num_test, dtype = self.ytr.dtype)

        # loop over all test rows
        for i in xrange(num_test):
            # find the nearest training image to the i'th test image
            # using the L1 distance (sum of absolute value differences)
            distances = np.sum(np.abs(self.Xtr - X[i,:]), axis = 1)
            min_index = np.argmin(distances) # get the index with smallest distance
            Ypred[i] = self.ytr[min_index] # predict the label of the nearest example

        return Ypred

```

Nearest Neighbor Classifier

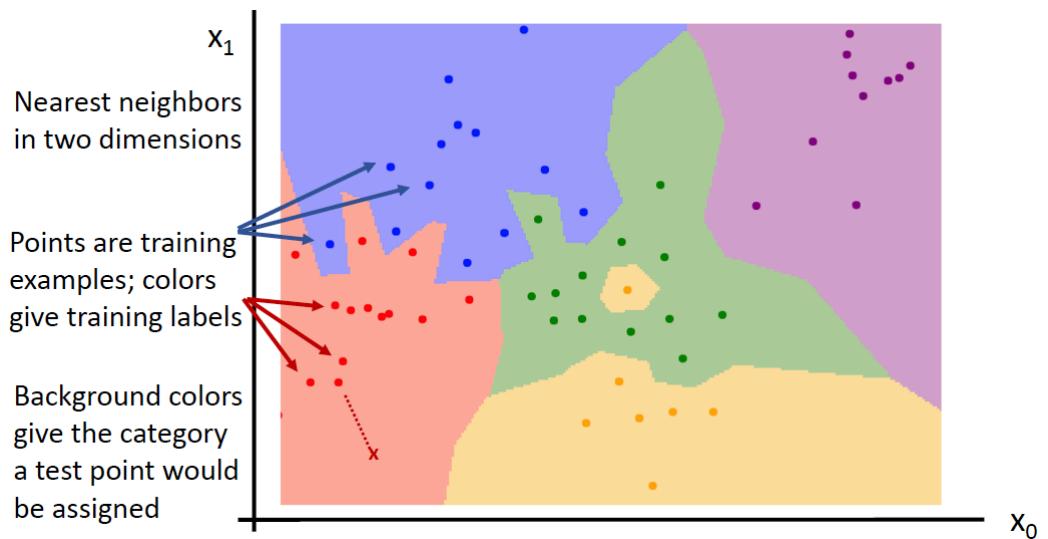
For each test image:
 Find nearest training image
 Return label of nearest image

- Cons:

- With N examples, training is O(1) while testing is O(N)
 --> bad: we can afford slow training, but need fast testing

- **Decision Boundary**

--> the boundary between two classification regions

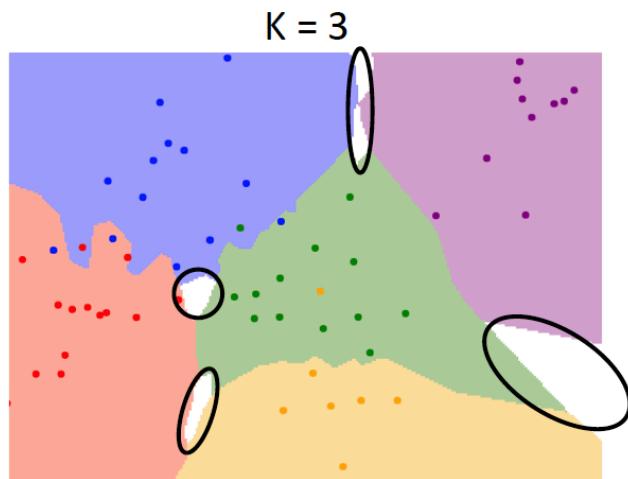


- Decision boundary can be noisy --> affected by outliers (yellow point in the graph)

- **K-Nearest Neighbors**

--> use more neighbors and average to **smooth out** desicion boundaries!

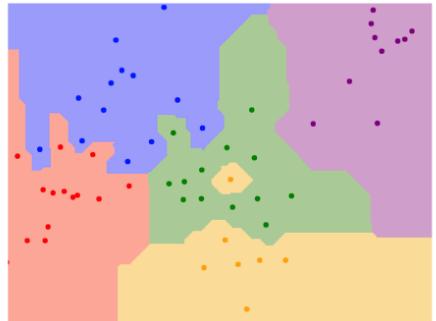
--> Instead of copying label from nearest neighbor, take majority vote from K closest points



- Distance Metric

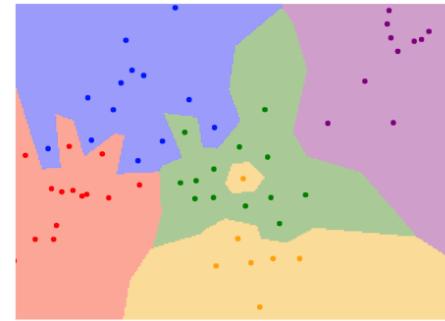
L1 (Manhattan) distance

$$d_1(I_1, I_2) = \sum_p |I_1^p - I_2^p|$$



L2 (Euclidean) distance

$$d_1(I_1, I_2) = \left(\sum_p (I_1^p - I_2^p)^2 \right)^{\frac{1}{2}}$$



- L1: all decision boundaries are $45^\circ/90^\circ/0^\circ$
- L2: decision boundary can be any angle
- With the right choice of distance metric, we can apply K-Nearest Neighbor to any type of data (image, similar articles...)
- KNN on raw pixels is seldom used
 - Very slow at test time
 - Distance metrics on pixels are not informative (not enough meaningful)



(all 3 images have same L2 distance to the one on the left)

image is

3. Hyperparameters

Hyperparameters: choices about our learning algorithm that we **don't learn** from the training data; instead we **set them at the start of the learning process**

Eg. distance metric, k...

--> problem-dependent: in general, need to try them all and see what works best for the data / task

- Setting Hyperparameters

Idea #1: Choose hyperparameters that work best on the data

BAD: K = 1 always works perfectly on training data

Your Dataset

Idea #2: Split data into **train** and **test**, choose hyperparameters that work best on test data

BAD: No idea how algorithm will perform on new data

train

test

Idea #3: Split data into **train**, **val**, and **test**; choose hyperparameters on val and evaluate on test

Better!

train

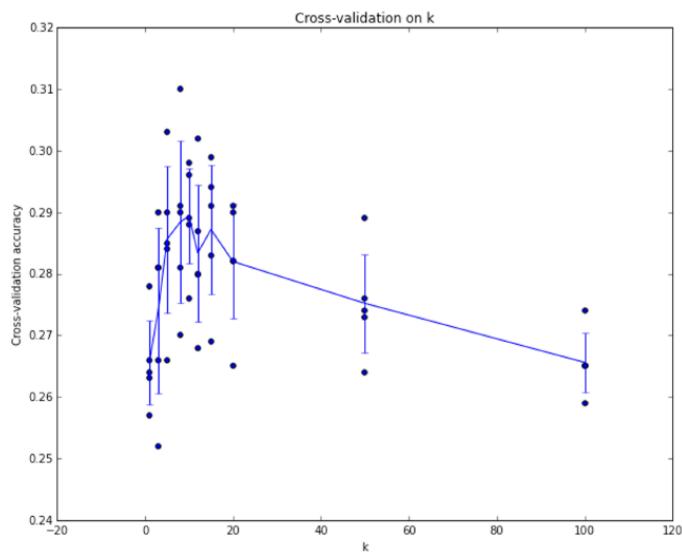
validation

test

Idea #4: Cross-Validation: Split data into **folds**, try each fold as validation and average the results

fold 1	fold 2	fold 3	fold 4	fold 5	test
fold 1	fold 2	fold 3	fold 4	fold 5	test
fold 1	fold 2	fold 3	fold 4	fold 5	test

Useful for small datasets, but (unfortunately) not used too frequently in deep learning



Example of 5-fold cross-validation for the value of k .

Each point: single outcome.

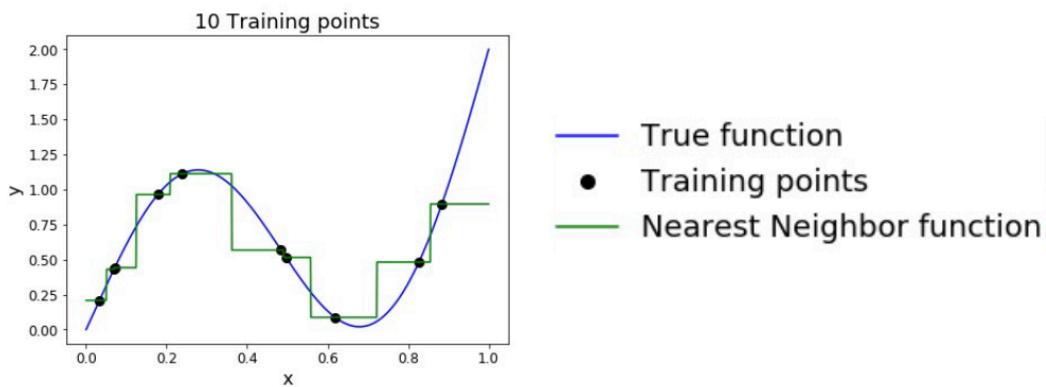
The line goes through the mean, bars indicated standard deviation

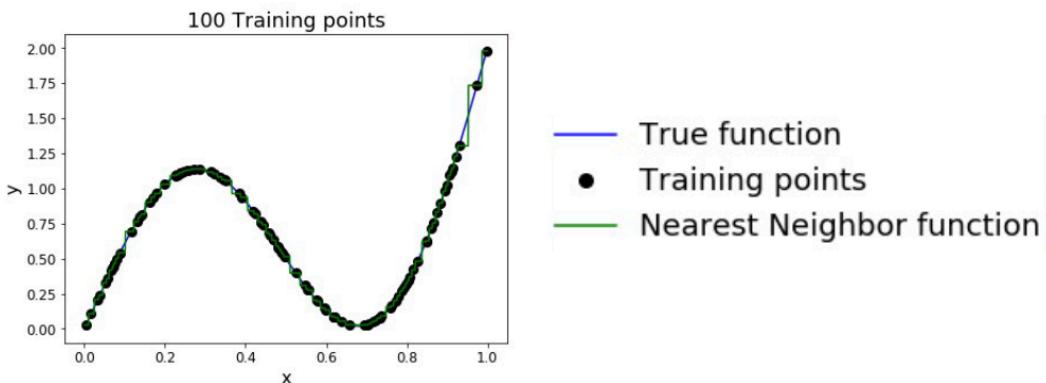
(Seems that $k \sim 7$ works best for this data)

Choose hyperparameters using the **validation set**, only run the test set at the very end!

4. Universal Approximation

As the number of training samples goes to infinity, nearest neighbor can *represent any function*!

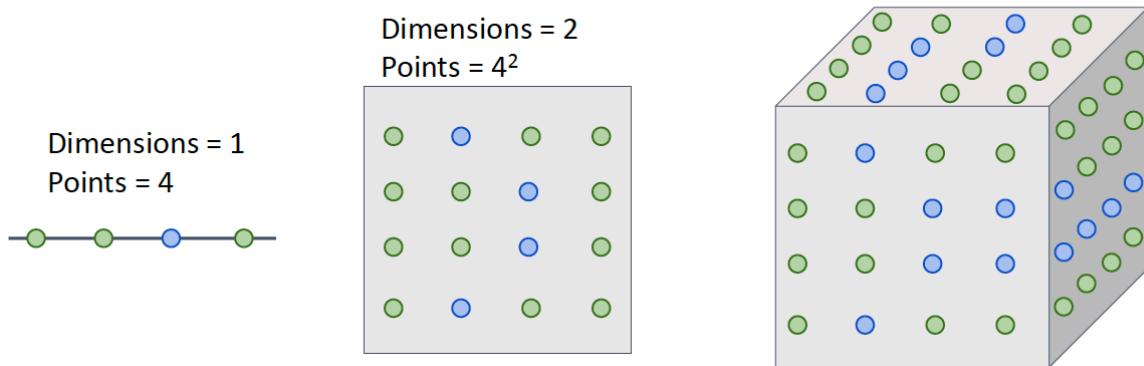




- Problem: **Curve of Dimensionality**

当数据在更高维时，训练数据点会指数型增长，无法进行训练

Curse of dimensionality: For uniform coverage of space, number of training points needed grows exponentially with dimension



eg. number of possible 32×32 binary images: $2^{32 \times 32} \approx 10^{308}$